

# PERFORMANCE ANALYSIS OF PARALLEL CODES ON HETEROGENEOUS SYSTEMS

---

E. Agullo, O. Aumage, B. Bramas, A. Buttari, A. Guermouche,  
F. Lopez, S. Nakov, S. Thibault

SOLHAR plenary meeting, Bordeaux 25-01-2026

## A MOTIVATING EXAMPLE

---

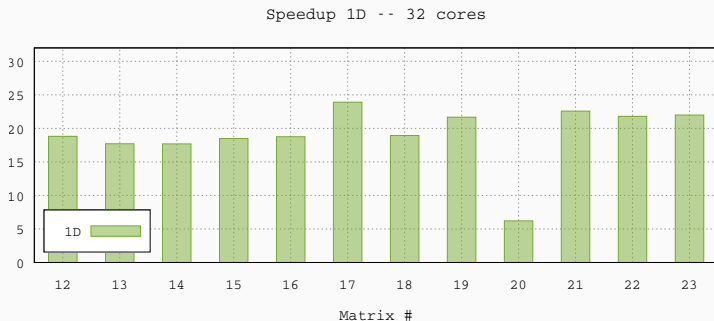
# Plain speedup is not enough

- `qr_mumps` + StarPU with 1D, block-column partitioning
- Matrices from UF collection

#	Matrix	Mflops			
12	hirlam	1384160	18	spa1_004	30335566
13	flower_8_4	2851508	19	n4c6-b6	62245957
14	Rucci1	5671282	20	sls	65607341
15	ch8-8-b3	10709211	21	TF18	194472820
16	GL7d24	16467844	22	lp_nug30	221644546
17	neos2	20170318	23	mk13-b5	259751609

- One node of the ADA supercomputer (IBM x3750-M4, Intel Sandy Bridge E5-4650 @ 2.7 GHz,  $4 \times 8$  cores)

# Experimental results: speedups



Speedup says something, e.g., performance is poor on small matrices and good on bigger ones.

Speedup doesn't say anything on **the reason**.

Is there a problem in the **implementation**, in the **algorithm**, in the **data**? what's that crappy matrix?

PERFORMANCE ANALYSIS APPROACH,  
HOMOGENEOUS CASE

---

THE

# Area performance upper bound

## Parallel efficiency

The **parallel efficiency** is defined as

$$e(p) = \frac{t^{\min}(p)}{t(p)} = \frac{\tilde{t}(1)}{t(p) \cdot p}$$

- $\tilde{t}(1)$  is the execution time of the **best sequential algorithm** on one core;
- $t(p)$  is the execution time of the **best parallel algorithm** on  $p$  cores.

Note that, in general,  $t(1) \geq \tilde{t}(1)$  because:

- parallelism requires partitioning of data and operations which reduces the efficiency of tasks;
- the parallel algorithm may trade some extra flops for concurrency.

# A finer performance analysis

The execution time  $t(p)$  can be decomposed in the following three terms:

- $t_t(p)$ : the time spent executing tasks.
- $t_r(p)$ : the overhead of the runtime system.  $t_r(1) := 0$ .
- $t_i(p)$ : idle time.  $t_i(1) := 0$ .

The overall efficiency can thus be written as:

$$\begin{aligned} e(p) &= \frac{\tilde{t}_t(1)}{t_t(p) + t_r(p) + t_i(p)} \\ &= \overbrace{\frac{\tilde{t}_t(1)}{t_t(1)}}^{e_g} \cdot \overbrace{\frac{t_t(1)}{t_t(p)}}^{e_t} \cdot \overbrace{\frac{t_t(p)}{t_t(p) + t_r(p)}}^{e_r} \cdot \overbrace{\frac{t_t(p) + t_r(p) + t_c(p)}{t_t(p) + t_r(p) + t_c(p) + t_i(p)}}^{e_p}. \end{aligned}$$

with:

# A finer performance analysis

The execution time  $t(p)$  can be decomposed in the following three terms:

- $t_t(p)$ : the time spent executing tasks.
- $t_r(p)$ : the overhead of the runtime system.  $t_r(1) := 0$ .
- $t_i(p)$ : idle time.  $t_i(1) := 0$ .

The overall efficiency can thus be written as:

$$\begin{aligned} e(p) &= \frac{\tilde{t}_t(1)}{t_t(p) + t_r(p) + t_i(p)} \\ &= \overbrace{\frac{\tilde{t}_t(1)}{t_t(1)}}^{e_g} \cdot \overbrace{\frac{t_t(1)}{t_t(p)}}^{e_t} \cdot \overbrace{\frac{t_t(p)}{t_t(p) + t_r(p)}}^{e_r} \cdot \overbrace{\frac{t_t(p) + t_r(p) + t_c(p)}{t_t(p) + t_r(p) + t_c(p) + t_i(p)}}^{e_p}. \end{aligned}$$

with:

$e_g$ : the **granularity efficiency**. Measures the impact exploiting of parallel algorithms compared to sequential ones.



# A finer performance analysis

The execution time  $t(p)$  can be decomposed in the following three terms:

- $t_t(p)$ : the time spent executing tasks.
- $t_r(p)$ : the overhead of the runtime system.  $t_r(1) := 0$ .
- $t_i(p)$ : idle time.  $t_i(1) := 0$ .

The overall efficiency can thus be written as:

$$\begin{aligned} e(p) &= \frac{\tilde{t}_t(1)}{t_t(p) + t_r(p) + t_i(p)} \\ &= \overbrace{\frac{\tilde{t}_t(1)}{t_t(1)}}^{e_g} \cdot \overbrace{\frac{t_t(1)}{t_t(p)}}^{e_t} \cdot \overbrace{\frac{t_t(p)}{t_t(p) + t_r(p)}}^{e_r} \cdot \overbrace{\frac{t_t(p) + t_r(p) + t_c(p)}{t_t(p) + t_r(p) + t_c(p) + t_i(p)}}^{e_p}. \end{aligned}$$

with:

$e_t$ : the **task efficiency**. Measures the exploitation of data locality.

# A finer performance analysis

The execution time  $t(p)$  can be decomposed in the following three terms:

- $t_t(p)$ : the time spent executing tasks.
- $t_r(p)$ : the overhead of the runtime system.  $t_r(1) := 0$ .
- $t_i(p)$ : idle time.  $t_i(1) := 0$ .

The overall efficiency can thus be written as:

$$\begin{aligned} e(p) &= \frac{\tilde{t}_t(1)}{t_t(p) + t_r(p) + t_i(p)} \\ &= \overbrace{\frac{\tilde{t}_t(1)}{t_t(1)}}^{e_g} \cdot \overbrace{\frac{t_t(1)}{t_t(p)}}^{e_t} \cdot \overbrace{\frac{t_t(p)}{t_t(p) + t_r(p)}}^{e_r} \cdot \overbrace{\frac{t_t(p) + t_r(p) + t_c(p)}{t_t(p) + t_r(p) + t_c(p) + t_i(p)}}^{e_p}. \end{aligned}$$

with:

$e_r$ : the **runtime efficiency**. Measures how the runtime overhead affects performance.

# A finer performance analysis

The execution time  $t(p)$  can be decomposed in the following three terms:

- $t_t(p)$ : the time spent executing tasks.
- $t_r(p)$ : the overhead of the runtime system.  $t_r(1) := 0$ .
- $t_i(p)$ : idle time.  $t_i(1) := 0$ .

The overall efficiency can thus be written as:

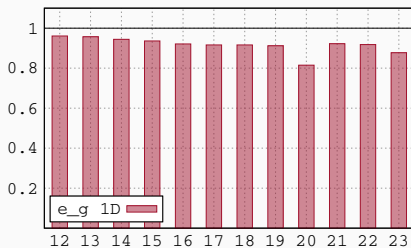
$$\begin{aligned} e(p) &= \frac{\tilde{t}_t(1)}{t_t(p) + t_r(p) + t_i(p)} \\ &= \overbrace{\frac{\tilde{t}_t(1)}{t_t(1)}}^{e_g} \cdot \overbrace{\frac{t_t(1)}{t_t(p)}}^{e_t} \cdot \overbrace{\frac{t_t(p)}{t_t(p) + t_r(p)}}^{e_r} \cdot \overbrace{\frac{t_t(p) + t_r(p) + t_c(p)}{t_t(p) + t_r(p) + t_c(p) + t_i(p)}}^{e_p}. \end{aligned}$$

with:

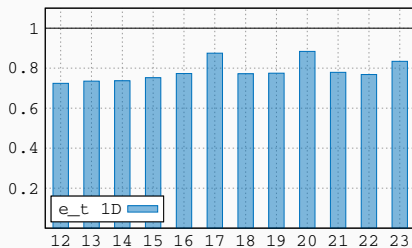
$e_p$ : the **pipeline efficiency**. Measures how much concurrency is available and how well it is exploited.

# Experimental results: efficiency breakdown

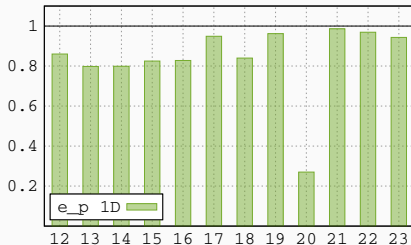
Granularity efficiency



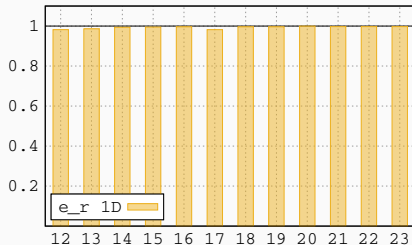
Task efficiency



Pipeline efficiency



Runtime efficiency



# 2D partitioning + CA front factorization

1D partitioning is not good for (strongly) **overdetermined** matrices:

- ▼ Most fronts are overdetermined
- ▲ The problem is mitigated by concurrent front factorizations

- 2D block partitioning (not necessarily square)
- Communication avoiding algorithms

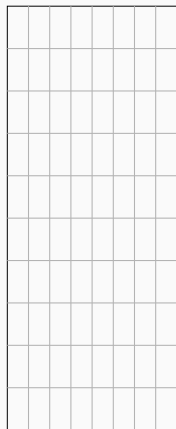
▲ More concurrency

▼ More complex dependencies

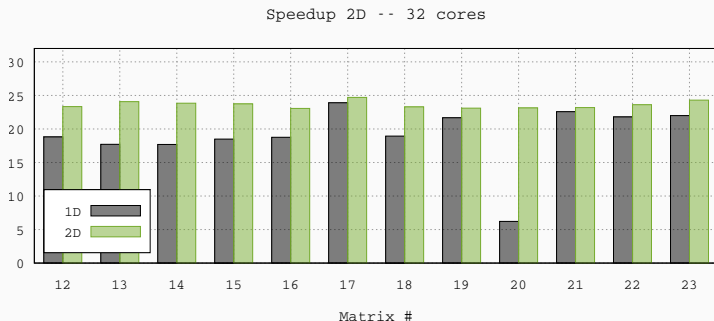
▼ Many more tasks (higher runtime overhead)

▼ Finer task granularity (less kernel efficiency)

Thanks to the simplicity of the STF programming model it is possible to plug in **2D methods** for factorizing the frontal matrices with a relatively moderate effort



# Experimental results: speedups

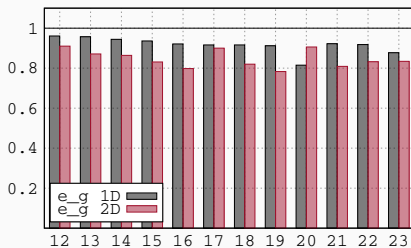


The scalability of the task-based multifrontal method is enhanced by the the introduction of 2D CA algorithms:

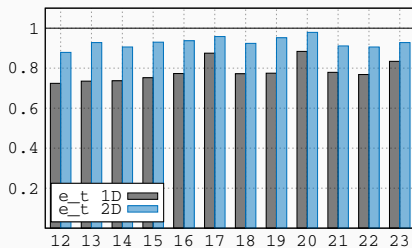
- Speedups are **uniform** for all tested matrices.
- We perform a comparative performance analysis wrt to the 1D case to show the benefits of the 2D scheme.

# Experimental results: efficiency breakdown

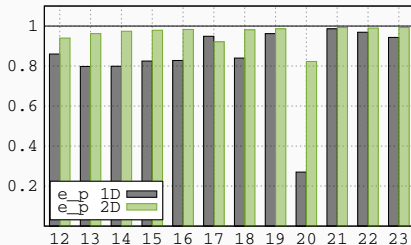
## Granularity efficiency



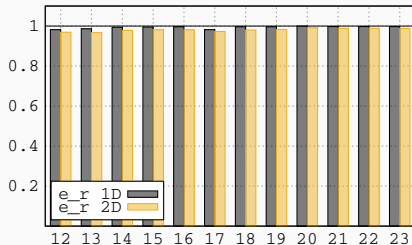
## Task efficiency



## Pipeline efficiency



## Runtime efficiency

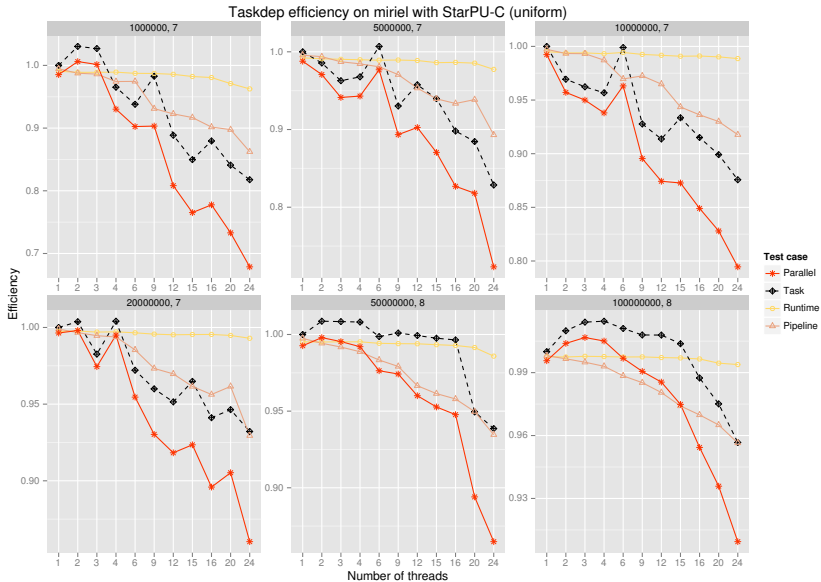


## CASE STUDY WITH SCALFMM

---

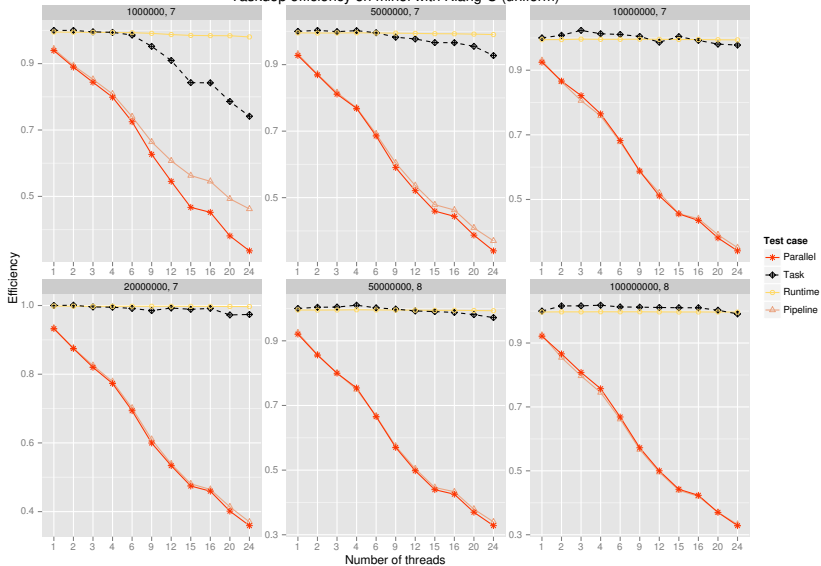


# Uniform - native StarPU (with commute)



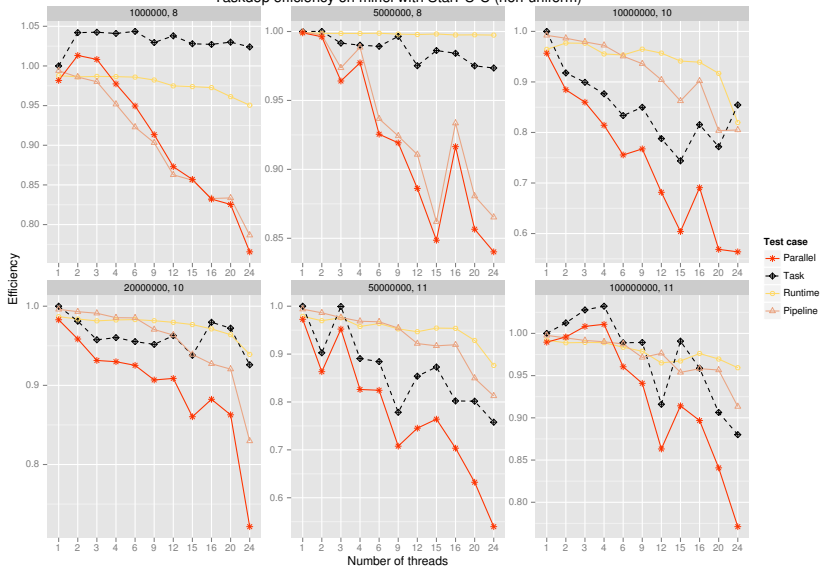
# Uniform - OpenMP-Klang-StarPU (with commute)

Taskdep efficiency on miriel with Klang-C (uniform)



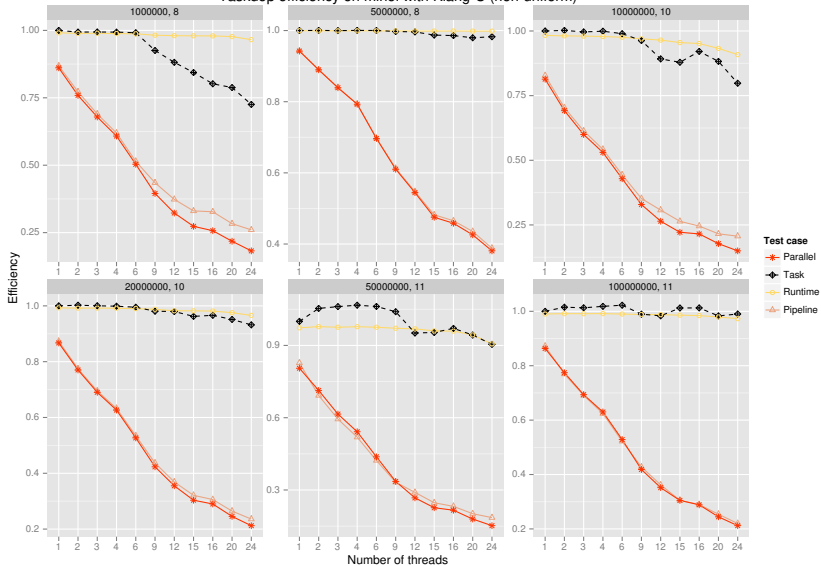
# Ellipsoid - native StarPU (with commute)

Taskdep efficiency on mrirel with StarPU-C (non-uniform)



# Ellipsoid - OpenMP-Klang-StarPU (with commute)

Taskdep efficiency on miriel with Klang-C (non-uniform)



PERFORMANCE ANALYSIS APPROACH,  
HETEROGENEOUS CASE

---


THE

# Area performance upper bound

The parallel efficiency can be defined as

$$e(p) = \frac{t^{min}(p)}{t(p)}$$

where  $t^{min}(p)$  is a lower bound on execution time on  $p$  resources corresponding to the **best schedule** under the following assumptions:




# Area performance upper bound

The parallel efficiency can be defined as

$$e(p) = \frac{t^{min}(p)}{t(p)}$$

where  $t^{min}(p)$  is a lower bound on execution time on  $p$  resources corresponding to the **best schedule** under the following assumptions:

1. No runtime overhead and no communications.




# Area performance upper bound

The parallel efficiency can be defined as

$$e(p) = \frac{t^{min}(p)}{t(p)}$$

where  $t^{min}(p)$  is a lower bound on execution time on  $p$  resources corresponding to the **best schedule** under the following assumptions:

1. No runtime overhead and no communications.
2. No tasks dependencies.





# Area performance upper bound

The parallel efficiency can be defined as

$$e(p) = \frac{t^{min}(p)}{t(p)}$$

where  $t^{min}(p)$  is a lower bound on execution time on  $p$  resources corresponding to the **best schedule** under the following assumptions:

1. No runtime overhead and no communications.
2. No tasks dependencies.
3. Tasks are moldable.




# Area performance upper bound

The parallel efficiency can be defined as

$$e(p) = \frac{t^{min}(p)}{t(p)}$$

where  $t^{min}(p)$  is a lower bound on execution time on  $p$  resources corresponding to the **best schedule** under the following assumptions:

1. No runtime overhead and no communications.
2. No tasks dependencies.
3. Tasks are moldable.



In the heterogeneous case we have  $t^{area}(p)$  is the solution of a **linear program**.


# Area performance upper bound

The parallel efficiency can be defined as

$$e(p) = \frac{t^{min}(p)}{t(p)}$$

where  $t^{min}(p)$  is a lower bound on execution time on  $p$  resources corresponding to the **best schedule** under the following assumptions:

1. No runtime overhead and no communications.
2. No tasks dependencies.
3. Tasks are moldable.



We consider  $t^{area}(p)$  when there is no performance loss resulting from the parallelization:  $\tilde{t}^{area}(p)$ .

# A finer performance analysis

The execution time  $t(p)$  can be decomposed in the following four terms:

- $t_t(p)$ : the time spent executing tasks.
- $t_r(p)$ : the overhead of the runtime system.
- $t_i(p)$ : idle time.

The overall efficiency can thus be written as:

$$\begin{aligned} e(p) &= \frac{\tilde{t}^{area}_t(p)}{t(p)} = \frac{\tilde{t}^{area}_t(p) \times p}{t_t(p) + t_r(p) + t_c(p) + t_i(p)} = \frac{\tilde{t}^{area}_t(p)}{t_t(p) + t_r(p) + t_c(p) + t_i(p)} \\ &= \frac{\overbrace{\tilde{t}^{area}_t(p)}^{e_g}}{t_t^{area}(p)} \cdot \frac{\overbrace{t_t^{area}(p)}^{e_t}}{t_t(p)} \cdot \frac{\overbrace{t_t(p)}^{e_r}}{t_t(p) + t_r(p)} \cdot \frac{\overbrace{t_t(p) + t_r(p) + t_c(p)}^{e_p}}{t_t(p) + t_r(p) + t_c(p) + t_i(p)}. \end{aligned}$$

with:

# A finer performance analysis

The execution time  $t(p)$  can be decomposed in the following four terms:

- $t_t(p)$ : the time spent executing tasks.
- $t_r(p)$ : the overhead of the runtime system.
- $t_i(p)$ : idle time.
- $t_c(p)$ : the time spent performing communications.

The overall efficiency can thus be written as:

$$\begin{aligned} e(p) &= \frac{\tilde{t}^{area}_t(p)}{t(p)} = \frac{\tilde{t}^{area}_t(p) \times p}{t_t(p) + t_r(p) + t_c(p) + t_i(p)} = \frac{\tilde{t}^{area}_t(p)}{t_t(p) + t_r(p) + t_c(p) + t_i(p)} \\ &= \overbrace{\frac{\tilde{t}^{area}_t(p)}{\tilde{t}^{area}_t(p)}}^{e_g} \cdot \overbrace{\frac{\tilde{t}^{area}_t(p)}{t_t(p)}}^{e_t} \cdot \overbrace{\frac{t_t(p)}{t_t(p) + t_r(p)}}^{e_r} \cdot \overbrace{\frac{t_t(p) + t_r(p)}{t_t(p) + t_r(p) + t_c(p)}}^{e_c} \cdot \overbrace{\frac{t_t(p) + t_r(p) + t_c(p)}{t_t(p) + t_r(p) + t_c(p) + t_i(p)}}^{e_p}. \end{aligned}$$

with:

$e_c$ : the **communication efficiency**. measures the cost of communications with respect to the actual work done due to data transfers between workers.

# A finer performance analysis

The execution time  $t(p)$  can be decomposed in the following four terms:

- $t_t(p)$ : the time spent executing tasks.
- $t_r(p)$ : the overhead of the runtime system.
- $t_i(p)$ : idle time.
- $t_c(p)$ : the time spent performing communications.

The overall efficiency can thus be written as:


$$\begin{aligned} e(p) &= \frac{\tilde{t}^{area}(p)}{t(p)} = \frac{\tilde{t}^{area}(p) \times p}{t_t(p) + t_r(p) + t_c(p) + t_i(p)} = \frac{\tilde{t}_t^{area}(p)}{t_t(p) + t_r(p) + t_c(p) + t_i(p)} \\ &= \overbrace{\frac{\tilde{t}_t^{area}(p)}{t_t^{area}(p)}}^{e_g} \cdot \overbrace{\frac{t_t^{area}(p)}{t_t(p)}}^{e_t} \cdot \overbrace{\frac{t_t(p)}{t_t(p) + t_r(p)}}^{e_r} \cdot \overbrace{\frac{t_t(p) + t_r(p)}{t_t(p) + t_r(p) + t_c(p)}}^{e_c} \cdot \overbrace{\frac{t_t(p) + t_r(p) + t_c(p)}{t_t(p) + t_r(p) + t_c(p) + t_i(p)}}^{e_p} \end{aligned}$$

with:

$e_t$ : the **task efficiency**. Measures how well the assignment of tasks to processing units matches the tasks properties to the units capabilities.


# Experimental results: absolute performance

- `qr_mumps` + StarPU with hierarchical partitioning and HeteroPrio++ scheduler
- One node of the Sirocco computer (Haswell Intel Xeon E5-2680 @ 2.5 GHz,  $2 \times 12$  cores + Nvidia K40)




# Experimental results: efficiency breakdown


Granularity efficiency




Task efficiency



Pipeline efficiency





Runtime efficiency






# Experimental results: efficiency breakdown



## CRITICAL PATH ANALYSIS

---

# Critical path analysis



$$\text{max\_speedup} = \text{avg\_concurrency} = \frac{\sum_{i \in \text{DAG}} w_i}{\sum_{i \in \text{CP}} w_i}$$

- The **DAG** used to conduct this analysis is the one related to the case where **32 working threads** are used.
- The **weight** of tasks is chosen to be equal to the **execution time** measured in an execution with only one working thread.

Thanks!

Questions?